


linux系統編程之信號（七）：被信號中斷的系統調用和庫函數處理方式

一些IO系統調用執行時, 如read 等待輸入期間, 如果收到一個信號, 系統將中斷read, 轉而執行信號處理函數. 當信號處理返回後, 系統遇到了一個問題: 是重新開始這個系統調用, 還是讓系統調用失敗? 早期UNIX系統的做法是, 中斷系統調用, 並讓系統調用失敗, 比如read返回-1, 同時設置errno 為EINTR中斷了的系統調用是沒有完成的調用, 它的失敗是臨時性的, 如果再次調用則可能成功, 這並不是真正的失敗, 所以要對這種情況進行處理, 典型的方式為:



```
while ( 1 ) {

    n = read(fd, buf, BUFSIZ);

    if (n == - 1 && errno != EINTR) {

        printf( " read error\n " );

        break ;

    }


    if (n == 0 ) {

        printf( " read done\n " );

        break ;


    }

}
```



這樣做邏輯比較繁瑣, 事實上, 我們可以從信號的角度來解決這個問題, 安裝信號的時候, 設置SA_RESTART屬性, 那麼當信號處理函數返回後, 被該信號中斷的系統調用將自動恢復.

示例程序：



```
#include <signal.h>
#include <stdio.h>
#include <stdlib.h>
#include <error.h>
#include <string .h>
#include <unistd.h>

void sig_handler( int signum)
{
    printf( " in handler\n " );
    sleep( 1 );
    printf( " handler return\n " );
}

int main( int argc, char ** argv)
```

```
{  
    char buf[ 100 ];  
    int ret;  
    struct sigaction action, old_action;  
  
    action.sa_handler = sig_handler;  
    sigemptyset( & action.sa_mask);  
    action.sa_flags = 0 ;  
    /* 版本1:不設置SA_RESTART屬性  
    *版本2:設置SA_RESTART屬性*/  
    // action.sa_flags |= SA_RESTART;  
  
    sigaction(SIGINT, NULL, & old_action);  
    if (old_action.sa_handler != SIG_IGN) {  
        sigaction(SIGINT, & action, NULL);  
    }  
  
    bzero(buf, 100 );  
  
    ret = read( 0 , buf, 100 );  
    if (ret == - 1 ) {  
        perror( " read " );  
    }  
  
    printf( " read %d bytes:\n " , ret);  
    printf( " %s\n " , buf);  
  
    return 0 ;  
}
```



當sa_flags不設置：SA_RESTART時：

結果：

```
[zxy@test unixenv_c]$ cc sa_restart.c  
[zxy@test unixenv_c]$ ./a.out  
^Cin handler  
handler return  
read: Interrupted system call  
read -1 bytes:  
  
[zxy@test unixenv_c]$ |
```

設置後：

當被中斷後，重新執行

```
[zxy@test unixenv_c]$ cc sa_restart.c
[zxy@test unixenv_c]$ ./a.out
^Cin handler
handler return
^Cin handler
handler return
hello world
read 12 bytes:
hello world

[zxy@test unixenv_c]$ |
```

man幫助說明：

Interruption of system calls and library functions by signal handlers

If a signal handler is invoked while a system call or library function call is blocked, then either:

- * the call is automatically restarted after the signal handler returns; or
- * the call fails with the error **EINTR**.

Which of these two behaviors occurs depends on the interface and whether or not the signal handler was established using the **SA_RESTART** flag (see [sigaction\(2\)](#)). The details vary across UNIX systems; below, the details for Linux.

If a blocked call to one of the following interfaces is interrupted by a signal handler, then the call will be automatically restarted after the signal handler returns if the **SA_RESTART** flag was used; otherwise the call will fail with the error **EINTR**:

- * [read\(2\)](#), [readv\(2\)](#), [write\(2\)](#), [writev\(2\)](#), and [ioctl\(2\)](#) calls on "slow" devices. A "slow" device is one where the I/O call may block for an indefinite time, for example, a terminal, pipe, or socket. (A disk is not a slow device according to this definition.) If an I/O call on a slow device has already transferred some data by the time it is interrupted by a signal handler, then the call will return a success status (normally, the number of bytes transferred).
- * [open\(2\)](#), if it can block (eg, when opening a FIFO; see [fifo\(7\)](#)).
- * [wait\(2\)](#), [wait3\(2\)](#), [wait4\(2\)](#), [waitid\(2\)](#), and [waitpid\(2\)](#).
- * Socket interfaces: [accept\(2\)](#), [connect\(2\)](#), [recv\(2\)](#), [recvfrom\(2\)](#), [recvmsg\(2\)](#), [send\(2\)](#), [sendto\(2\)](#), and [sendmsg\(2\)](#), unless a timeout has been set on the socket (see below).
- * File locking interfaces: [flock\(2\)](#) and [fcntl\(2\)](#) **F_SETLK**.

- * POSIX message queue interfaces: mq_receive(3) ,
mq_timedreceive(3) , mq_send(3) , and mq_timedsend(3) .
- * futex(2) **FUTEX_WAIT** (since Linux 2.6.22; beforehand, always failed with **EINTR**).
- * POSIX semaphore interfaces: sem_wait(3) and sem_timedwait(3) (since Linux 2.6.22; beforehand, always failed with **EINTR**).

The following interfaces are never restarted after being interrupted by a signal handler, regardless of the use of **SA_RESTART** ; they always fail with the error **EINTR** when interrupted by a signal handler:

- * Socket interfaces, when a timeout has been set on the socket using setsockopt(2) : accept(2) , recv(2) , recvfrom(2) , and recvmsg(2) , if a receive timeout (**SO_RCVTIMEO**) has been set; connect(2) , send(2) , sendto(2) , and sendmsg(2) , if a send timeout (**SO_SNDTIMEO**) has been set.
- * Interfaces used to wait for signals: pause(2) , sigsuspend(2) , sigtimedwait(2) , and sigwaitinfo(2) .
- * File descriptor multiplexing interfaces: epoll_wait(2) , epoll_pwait(2) , poll(2) , ppoll(2) , select(2) , and pselect(2) .
- * System V IPC interfaces: msgrcv(2) , msgsnd(2) , semop(2) , and semtimedop(2) .
- * Sleep interfaces: clock_nanosleep(2) , nanosleep(2) , and usleep(3) .
- * read(2) from an inotify(7) file descriptor.
- * io_getevents(2) .

The sleep(3) function is also never restarted if interrupted by a handler, but gives a success return: the number of seconds remaining to sleep.
