

linux系統編程之管道（二）：管道讀寫規則

一，管道讀寫規則

當沒有數據可讀時

- O_NONBLOCK disable：read調用阻塞，即進程暫停執行，一直等到有數據來到為止。
- O_NONBLOCK enable：read調用返回-1，errno值為EAGAIN。

當管道滿的時候

- O_NONBLOCK disable：write調用阻塞，直到有進程讀走數據
- O_NONBLOCK enable：調用返回-1，errno值為EAGAIN

如果所有管道寫端對應的文件描述符被關閉，則read返回0

如果所有管道讀端對應的文件描述符被關閉，則write操作會產生信號SIGPIPE

當要寫入的數據量不大於PIPE_BUF時，linux將保證寫入的原子性。

當要寫入的數據量大於PIPE_BUF時，linux將不再保證寫入的原子性。

二，驗證示例

示例一：O_NONBLOCK disable：read調用阻塞，即進程暫停執行，一直等到有數據來到為止。



```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <fcntl.h>

int main( void )
{
    int fds[ 2 ];
    if (pipe(fds) == - 1 ){
        perror( " pipe error " );
        exit(EXIT_FAILURE);
    }
    pid_t pid;
    pid = fork();
    if (pid == - 1 ){
        perror( " fork error " );
        exit(EXIT_FAILURE);
    }
    if (pid == 0 ){
        close(fds[ 0 ]); // 子進程關閉讀端
        sleep( 10 );
        write(fds[ 1 ], " hello " , 5 );
        exit(EXIT_SUCCESS);
    }

    close(fds[ 1 ]); // 父進程關閉寫端
    char buf[ 10 ] = { 0 };
    read(fds[ 0 ],buf, 10 );
```

```
printf( " receive datas = %s\n " ,buf);  
return 0 ;  
}
```



結果：

```
[zxy@test unixenv_c]$ cc pipe03.c  
[zxy@test unixenv_c]$ ./a.out
```

說明：管道創建時默認打開了文件描述符，且默認是阻塞（block）模式打開

所以這裡，我們讓子進程先睡眠10s，父進程因為沒有數據從管道中讀出，被阻塞了，直到子進程睡眠結束，向管道中寫入數據後，父進程才讀到數據

示例二：O_NONBLOCK enable：read調用返回-1，errno值為EAGAIN。



```
#include <stdio.h>  
#include <unistd.h>  
#include <stdlib.h>  
#include <fcntl.h>  
  
int main( void )  
{  
    int fds[ 2 ];  
    if (pipe(fds) == - 1 ){  
        perror( " pipe error " );  
        exit(EXIT_FAILURE);  
    }  
    pid_t pid;  
    pid = fork();  
    if (pid == - 1 ){  
        perror( " fork error " );  
        exit(EXIT_FAILURE);  
    }  
    if (pid == 0 ){  
        close(fds[ 0 ]); // 子進程關閉讀端  
        sleep( 10 );  
        write(fds[ 1 ], " hello " , 5 );  
        exit(EXIT_SUCCESS);  
    }  
  
    close(fds[ 1 ]); // 父進程關閉寫端  
    char buf[ 10 ] = { 0 };  
    int flags = fcntl(fds[ 0 ], F_GETFL); // 先獲取原先的flags  
    fcntl(fds[ 0 ],F_SETFL,flags | O_NONBLOCK); // 設置fd為阻塞模式  
    int ret;  
    ret = read(fds[ 0 ],buf, 10 );  
    if (ret == - 1 ){  
  
        perror( " read error " );  
    }  
}
```

```
        exit(EXIT_FAILURE);
    }

    printf( " receive datas = %s\n " ,buf);
    return 0 ;
}
```



結果：

```
[zxy@test unixenv_c]$ cc pipe04.c
[zxy@test unixenv_c]$ ./a.out
read error: Resource temporarily unavailable
[zxy@test unixenv_c]$ |
```

示例三：如果所有管道寫端對應的文件描述符被關閉，則read返回0



```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <fcntl.h>

int main( void )
{
    int fds[ 2 ];
    if (pipe(fds) == - 1 ){
        perror( " pipe error " );
        exit(EXIT_FAILURE);
    }
    pid_t pid;
    pid = fork();
    if (pid == - 1 ){
        perror( " fork error " );
        exit(EXIT_FAILURE);
    }
    if (pid == 0 ){
        close(fds[ 1 ]); // 子進程關閉寫端
        exit(EXIT_SUCCESS);
    }

    close(fds[ 1 ]); // 父進程關閉寫端
    char buf[ 10 ] = { 0 };

    int ret;
    ret = read(fds[ 0 ],buf, 10 );
    printf( " ret = %d\n " , ret);

    return 0 ;
}
```



結果：

```
[zxy@test unixenv_c]$ cc pipe05.c
[zxy@test unixenv_c]$ ./a.out
ret = 0
[zxy@test unixenv_c]$ |
```

可知確實返回0，表示讀到了文件末尾，並不表示出錯

示例四：如果所有管道讀端對應的文件描述符被關閉，則write操作會產生信號SIGPIPE



```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <fcntl.h>
#include <signal.h>

void sighandler( int signo);
int main( void )
{
    int fds[ 2 ];
    if (signal(SIGPIPE,sighandler) == SIG_ERR)
    {
        perror( " signal error " );
        exit(EXIT_FAILURE);
    }
    if (pipe(fds) == - 1 ){
        perror( " pipe error " );
        exit(EXIT_FAILURE);
    }
    pid_t pid;
    pid = fork();
    if (pid == - 1 ){
        perror( " fork error " );
        exit(EXIT_FAILURE);
    }
    if (pid == 0 ){
        close(fds[ 0 ]); // 子進程關閉讀端
        exit(EXIT_SUCCESS);
    }

    close(fds[ 0 ]); // 父進程關閉讀端
    sleep( 1 ); // 確保子進程也將讀端關閉
    int ret;
    ret = write(fds[ 1 ], " hello " , 5 );
    if (ret == - 1 ){
        printf( " write error\n " );
    }
    return 0 ;
}

void sighandler( int signo)
{
    printf( " catch a SIGPIPE signal and signum = %d\n " ,signo);
}
```



結果：

```
[zxy@test unixenv_c]$ cc pipe06.c
[zxy@test unixenv_c]$ ./a.out
catch a SIGPIPE signal and signum = 13
write error
[zxy@test unixenv_c]$ |
```

可知當所有讀端都關閉時，write時確實產生SIGPIPE信號

示例五：O_NONBLOCK disable：write調用阻塞，直到有進程讀走數據



```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <fcntl.h>

int main( void )
{
    int fds[ 2 ];
    if (pipe(fds) == - 1 ){
        perror( " pipe error " );
        exit(EXIT_FAILURE);
    }
    int ret;
    int count = 0 ;
    while ( 1 ){
        ret = write(fds[ 1 ], " A " , 1 ); // fds[1]默認是阻塞模式
        if (ret == - 1 ){
            perror( " write error " );
            break ;
        }
        count ++ ;
    }

    return 0 ;
}
```




結果：

```
[zxy@test unixenv_c]$ cc pipe07.c
[zxy@test unixenv_c]$ ./a.out
```

說明：fd打開時默認是阻塞模式，當pipe緩衝區滿時，write操作確實阻塞了，等待其他進程將數據從管道中取走


示例六：O_NONBLOCK enable：調用返回-1，errno值為EAGAIN



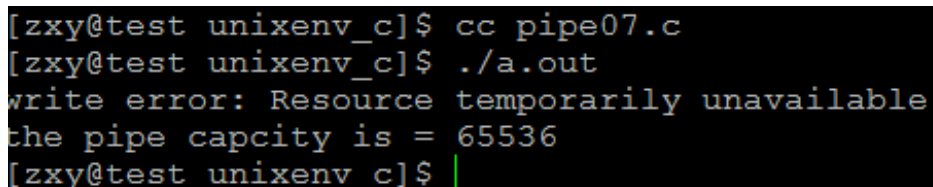
```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <fcntl.h>

int main( void )
{
    int fds[ 2 ];
    if (pipe(fds) == - 1 ){
        perror( " pipe error " );
        exit(EXIT_FAILURE);
    }
    int ret;
    int count = 0 ;
    int flags = fcntl(fds[ 1 ],F_GETFL);
    fcntl(fds[ 1 ],F_SETFL,flags| O_NONBLOCK);
    while ( 1 ){
        ret = write(fds[ 1 ], " A " , 1 ); // fds[1]默認是阻塞模式
        if (ret == - 1 ){
            perror( " write error " );
            break ;
        }
        count ++ ;
    }
    printf( " the pipe capacity is = %d\n " ,count);

    return 0 ;
}
```



結果：



```
[zxy@test unixenv_c]$ cc pipe07.c
[zxy@test unixenv_c]$ ./a.out
write error: Resource temporarily unavailable
the pipe capacity is = 65536
[zxy@test unixenv_c]$
```

可知也出現EAGAIN錯誤，管道容量是65536字節

man 7 pipe說明：

Pipe capacity

A pipe has a limited capacity. If the pipe is full, then a `write(2)` will block or fail, depending on whether the `O_NONBLOCK` flag is set (see below). Different implementations have different limits for the pipe capacity. Applications should not rely on a particular capacity: an application should be designed so that a reading process consumes data as soon as it is available, so that a writing process does not remain blocked.

In Linux versions before 2.6.11, the capacity of a pipe was the same

as the system page size (eg, 4096 bytes on i386). Since Linux 2.6.11, **the pipe capacity is 65536 bytes.**

三，管道寫與PIPE_BUF關係

man幫助說明：

PIPE_BUF

POSIX.1-2001 says that write(2) s of **less than PIPE_BUF bytes must be atomic: the output data is written to the pipe as a contiguous sequence.** Writes of **more than PIPE_BUF bytes may be nonatomic: the kernel may interleave the data with data written by other processes .**

POSIX.1-2001 requires **PIPE_BUF** to be at least 512 bytes. (On Linux, **PIPE_BUF** is 4096 bytes.) The precise semantics depend on whether the file descriptor is nonblocking (**O_NONBLOCK**), whether there are multiple writers to the pipe, and on n , the number of bytes to be written:

O_NONBLOCK disabled, $n \leq \text{PIPE_BUF}$

All n bytes are written atomically; write(2) may block if there is not room for n bytes to be written immediately

阻塞模式時且 $n < \text{PIPE_BUF}$ ：寫入具有原子性，如果沒有足夠的空間供 n 個字節全部寫入，則阻塞直到有足夠空間將 n 個字節全部寫入管道

O_NONBLOCK enabled, $n \leq \text{PIPE_BUF}$

If there is room to write n bytes to the pipe, then write(2) succeeds immediately, writing all n bytes; otherwise write(2) fails, with errno set to **EAGAIN** .

非阻塞模式時且 $n < \text{PIPE_BUF}$ ：寫入具有原子性，立即全部成功寫入，否則一個都不寫入，返回錯誤

O_NONBLOCK disabled, $n > \text{PIPE_BUF}$

The write is nonatomic: the data given to write(2) may be interleaved with write(2) s by other process; the write(2) blocks until n bytes have been written.

阻塞模式時且 $n > \text{PIPE_BUF}$ ：不具有原子性，可能中間有其他進程穿插寫入，直到將 n 字節全部寫入才返回，否則阻塞等待寫入

O_NONBLOCK enabled, $n > \text{PIPE_BUF}$

If the pipe is full, then write(2) fails, with errno set to **EAGAIN** . Otherwise, from 1 to n bytes may be written (ie, a "partial write" may occur; the caller should check the return value from write(2) to see how many bytes were actually written), and these bytes may be interleaved with writes by other processes.

非阻塞模式時且 $n > \text{PIPE_BUF}$ ：如果管道滿的，則立即失敗，一個都不寫入，返回錯誤，如果不滿，則返回寫入的字節數為 $1 \sim n$ ，即部分寫入，寫入時可能有其他進程穿插寫入

- 當要寫入的數據量不大於 PIPE_BUF 時，linux將保證寫入的原子性。
- 當要寫入的數據量大於 PIPE_BUF 時，linux將不再保證寫入的原子性。

注：管道容量不一定等於 PIPE_BUF

示例：當寫入數據大於 PIPE_BUF 時



```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
```

```
#include <unistd.h>
#include <sys/types.h>
#include <errno.h>
#include <fcntl.h >

#define ERR_EXIT(m) \
    do \
    { \
        perror(m); \
        exit(EXIT_FAILURE); \
    } while ( 0 )

#define TEST_SIZE 68*1024

int main( void )
{
    char a[TEST_SIZE];
    char b[TEST_SIZE];
    char c[TEST_SIZE];

    memset(a, ' A ', sizeof (a));
    memset(b, ' B ', sizeof (b));
    memset(c, ' C ', sizeof (c));

    int pipefd[ 2 ];

    int ret = pipe(pipefd);
    if (ret == - 1 )
        ERR_EXIT( " pipe error " );

    pid_t pid;
    pid = fork();
    if (pid == 0 ) // 第一個子進程
    {
        close(pipefd[ 0 ]);
        ret = write(pipefd[ 1 ], a, sizeof (a));
        printf( " apid=%d write %d bytes to pipe\n " , getpid(), ret);
        exit( 0 );
    }

    pid = fork();

    if (pid == 0 ) // 第二個子進程
    {
        close(pipefd[ 0 ]);
        ret = write(pipefd[ 1 ], b, sizeof (b));
        printf( " bpid=%d write %d bytes to pipe\n " , getpid(), ret);
        exit( 0 );
    }

    pid = fork();
```



```
if (pid == 0 ) // 第三個子進程
{
    close(pipefd[ 0 ]);
    ret = write(pipefd[ 1 ], c, sizeof (c));
    printf( " bpid=%d write %d bytes to pipe\n " , getpid(), ret);
    exit( 0 );
}

close(pipefd[ 1 ]);

sleep( 1 );
int fd = open( " test.txt " , O_WRONLY | O_CREAT | O_TRUNC, 0644 );
char buf[ 1024 * 4 ] = { 0 };
int n = 1 ;
while ( 1 )
{
    ret = read(pipefd[ 0 ], buf, sizeof (buf));
    if (ret == 0 )
        break ;
    printf( " n=%02d pid=%d read %d bytes from pipe buf[4095]=%c\n " , n++, getpid(), ret,
buf[ 4095 ]);
    write(fd, buf, ret);
}
return 0 ;
}
```



結果：

```
[zxy@test unixenv_c]$ cc pipe08.c
[zxy@test unixenv_c]$ ./a.out
n=01 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=02 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=03 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=04 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=05 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=06 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=07 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=08 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=09 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=10 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=11 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=12 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=13 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=14 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=15 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=16 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=17 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=18 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=19 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=20 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=21 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=22 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=23 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=24 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=25 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=26 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=27 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=28 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=29 pid=10082 read 4096 bytes from pipe buf[4095]=B
```

```
bpid=10085 write 69632 bytes to pipe
bpid=10084 write 69632 bytes to pipe
n=30 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=31 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=32 pid=10082 read 4096 bytes from pipe buf[4095]=B
n=33 pid=10082 read 4096 bytes from pipe buf[4095]=C
n=34 pid=10082 read 4096 bytes from pipe buf[4095]=B
apid=10083 write 69632 bytes to pipe
n=35 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=36 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=37 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=38 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=39 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=40 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=41 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=42 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=43 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=44 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=45 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=46 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=47 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=48 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=49 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=50 pid=10082 read 4096 bytes from pipe buf[4095]=A
n=51 pid=10082 read 4096 bytes from pipe buf[4095]=A
[zxy@test unixenv_c]$ |
```

可見各子進程間出現穿插寫入，並沒保證原子性寫入，且父進程在子進程編寫時邊讀。